

Exploring the Data Structures and Algorithms behind Version Control Systems

(Student Proposed)

Reece Donovan

Final Year Project
B.Sc. Computer Science

Supervisor: Prof. Ken Brown
Second Reader: Dr. Klaas-Jan Stol

Department of Computer Science
University College Cork

April 2023

Abstract

Declaration of Originality

Acknowledgements

Contents

1	Introduction	1
1.1	Motivation	2
1.2	Objective	2
2	Background	3
2.1	What is a Version Control System	3
2.2	Evolution of Version Control Systems	4
2.2.1	Local Version Control Systems (LVCS)	4
	Source Code Control System (SCCS)	5
	Revision Control System (RCS)	7
2.2.2	Centralized Version Control Systems (CVCS)	8
	Concurrent Versions System (CVS)	9
	Perforce Helix Core	10
	Apache Subversion (SVN)	11
2.2.3	Distributed Version Control Systems (DVCS)	12
	Git	13
	Mercurial	15
3	Design	16
3.1	Data Structures	16
3.1.1	Linked List	17
3.1.2	Binary Tree	20
3.1.3	Hash Table	22
3.1.4	Directed Acyclic Graph (DAG)	24
3.2	Algorithms	27
3.2.1	Searching	27
	Linear Search	27
	Binary Search	28
	Depth First Search (DFS)	29
	Breadth First Search (BFS)	30

3.2.2	Diffing Algorithms	32
	Myers Diff Algorithm	32
	Patience Diff Algorithm	33
4	Implementation	35
4.1	Data Structures	35
4.1.1	Revision Type	35
	Tests	36
	Benchmarks	37
4.1.2	Doubly Linked List	37
	Tests	39
	Benchmarks	40
4.1.3	Binary Search Tree	42
	Tests	44
	Benchmarks	44
4.1.4	Directed Acyclic Graph	46
	Tests	47
	Benchmarks	48
4.2	Search Algorithms	49
4.2.1	Linear Search	49
	Tests	50
	Benchmarks	50
4.2.2	Binary Search	52
	Tests	53
	Benchmarks	54
4.2.3	Depth-First Search	55
	Tests	57
	Benchmarks	57
4.2.4	Breadth-First Search	59
	Tests	60
	Benchmarks	60
4.3	Dynamic Programming Algorithms	62
4.3.1	Longest Common Subsequence	62
	Inefficient Algorithm	62
	Efficient Algorithm	62
	Tests	62
	Benchmarks	62
4.3.2	Longest Increasing Subsequence	62
	Inefficient Algorithm	62
	Efficient Algorithm	62

	Tests	62
	Benchmarks	62
4.4	Diffing Algorithms	62
4.4.1	Myers Diff	62
	Tests	62
	Benchmarks	62
4.4.2	Patience Diff	62
	Tests	62
	Benchmarks	62
5	Evaluation	63
6	Conclusion	64
	List of Figures	65
	List of Tables	66
	Bibliography	68

Chapter 1

Introduction

Version Control is a critical aspect of software development that helps developers keep track of changes made to a codebase. It enables software teams to work collaboratively, efficiently, and accurately on projects, reducing the likelihood of errors and conflicts. In essence, version control is the process of tracking and managing changes to files over time.

Version Control System (VCS) is a software tool that automates the version control process. It provides a centralised repository where developers can store their code, track changes, and collaborate with other team members. The Version Control System ensures that each team member has access to the latest version of the code and can work on it simultaneously.

The impact of **Version Control Systems** on software development has been immense. Before the advent of Version Control Systems, developers used to rely on manual processes to track changes, which was time-consuming and error-prone. With VCS, software teams can work together more efficiently, manage changes more effectively, and deliver better-quality software products.

Version Control Systems have also facilitated the rapid growth of **Continuous Integration** and **Continuous Delivery (CI/CD)**, which have become essential systems in modern large scale software development. Overall, Version Control Systems have revolutionised how software is developed, making it easier, faster, and more reliable.

1.1 Motivation

Over the years, Version Control Systems have become essential for software developers, enabling them to collaborate and work more effectively. As the software development industry has evolved, so have Version Control Systems, leading to the creation of many different software solutions, each with unique features and strengths. However, these systems' core data structures and algorithms have remained relatively unchanged.

The performance of Version Control Systems can be crucial in determining the efficiency of software development projects. Slow or inefficient Version Control Systems can result in delays, errors, and even project failure. Therefore, it is important to explore alternative data structures and algorithms that could improve the performance of these systems.

1.2 Objective

Version Control Systems (VCS) have become essential for software developers to collaborate effectively and efficiently. However, many users may not fully comprehend the intricacies of the underlying concepts at the core of these systems.

This report aimed to address this knowledge gap by exploring the underlying data structures and algorithms used to power Version Control Systems and by evaluating potential trade-offs and benefits of alternative approaches. In order to provide insight into how these factors impact system performance, scalability, and overall effectiveness.

In addition, this report aimed to provide a comprehensive overview of the evolution of Version Control Systems, including an overview of the most popular version control solutions throughout the years, such as Source Code Control System (SCCS)—1972, Apache Subversion (SVN)—2000, and Git—2005.

Chapter 2

Background

2.1 What is a Version Control System

A Version Control System (VCS) saves modifications made by individual software developers, allowing them to track their work over time and making the process more accessible. Furthermore, it facilitates sharing data between nodes, where each node can be kept up to date with the latest versions, minimising the need to handle merge conflicts.

Version Control Systems provide several advantages, such as aiding in collaboration among programmers and improving efficiency when working on large groups of mixed products. In addition, VCSs offer an audit trail and easy branching functionality, simplifying team-based development. It also enables an understanding of who has worked on specific pieces or sections during a given period.

By ensuring all changes are appropriately tracked, VCSs promote transparency within teams while maintaining optimal performance levels through streamlined management processes. This helps avoid errors arising from inconsistencies between versions, such as mismatched documents or programming problems resulting from conflicting edit operations.

Significance of Version Control Systems

For most software projects, the source code is a valuable asset that must be protected. Additionally, for many software teams, the source code represents a repository of invaluable knowledge and understanding about the problem that developers have amassed and refined through meticulous effort.

Version control safeguards this precious resource by preventing catastrophes, casual degradation, human error, or unintended consequences from affecting its quality. Software developers working in teams continually write new source code and modify

existing code, which is organised into file folders called "file trees". Developers may work on different tasks simultaneously, such as implementing new features or fixing unrelated bugs, by altering parts of the file tree.

Version control assists teams in addressing these challenges by tracking each change made by contributors, helping to prevent concurrent work from conflicting with one another. Consequently, any incompatibility can be detected and resolved without impeding team members' progress further down the line, ensuring that changes do not introduce bugs.

2.2 Evolution of Version Control Systems




Figure 2.1: Timeline of the Creation of Version Control Systems [9]

2.2.1 Local Version Control Systems (LVCS)

Local Version Control Systems are systems designed to operate on a single machine, early VCSs were intended to track changes for individual files, and checked-out files could only be edited locally by one user at a time [9]. In addition, they were built on the assumption that all users would log into the same shared Unix host with their own accounts, which was not always possible.

The main benefit of LVCS is that it provides a basic level of version control functionality without requiring any network connectivity. However, this can also be seen as a drawback when you consider that with LVCS, there is no central repository to store the code, which means that each developer has their own copy. This can lead to issues with merging changes and tracking changes across multiple copies.




Figure 2.2: Local Version Control System (LVCS)

Source Code Control System (SCCS)

Source Code Control System was released in 1972 and is one of the first successful VCS tools [9]. It was written by Marc Rochkind at Bell Labs, who wanted to solve the problem of tracking file revisions. The tool made tracking down bugs introduced into a program significantly more manageable. SCCS is worth understanding at a basic level because it helped set up modern VCS tools that developers use today.

Architecture

Much like modern VCS tools, SCCS has a set of commands that allow developers to work with the versioning of files. The basic command functionality is:

- Check-in files to track their history.
- Check-out specific file versions for review.

- Check-out specific file versions for editing.
- Check-in new file versions with comments explaining the changes.
- Revert changes made to a checked-out file.
- Basic branching and merging of changes.
- Print a log of a file's version history.

A particular type of file called an **s-file** or a **history file** is created when a file is tracked by SCCS. This file is named with the original filename prefixed with an **s.** and is stored in a subdirectory called **SCCS**.

So a file called **test.txt** would get a history file created in the **./SCCS/** directory with the name of **s.test.txt**. When created, the **s-file** contains the original file contents, a header that contains the file's version number, and some other metadata. There are also checksums stored in the **s-file** that are used to verify the integrity of the file (i.e. to ensure that the file has not been tampered with). The **s-file** content is not encoded or compressed in any way, which is a clear difference from modern VCS tools.

Since the original file's content is now stored in the history file, it can be retrieved into the working directory for review, compilation, or editing. Further changes made to this new copy, such as line additions, modifications and removals, can be checked back into a revised version of the history file, which increments its revision number [9].

Subsequent SCCS check-ins only store only deltas or changes to a file instead of storing entire contents each time; when a check-in is made, subsequent revisions are added onto existing delta tables inside an amended history file (history files do not use compression). This decreases the size of these large histories since they are not using compression on their files, so they take up more space than just having one complete copy that you are tracking with no differences like Word docs etc.

As previously mentioned, SCCS uses the Delta method known as Interleaved Deltas, which allows constant-time checkout regardless of how old your checked-out revision is - i.e., older revisions do not take longer than newer ones.

It is important to note that all files are tracked and checked in separately; there is no way to check in changes on multiple files as a part of one atomic unit - like a commit in Git. Each tracked file has its corresponding history file, which stores revisions. Generally, this means that the version numbers of different files cannot match each other. However, matching revision numbers can be achieved by editing every file at once (even if not all of the changed areas have fundamental changes) and

checking them all together. This will increment revision numbers for any modified content, so they will be consistent with their revisions—but it’s not comparable to including lots of different pieces within a single commit like you would in Git. In SCCS, these make individual check-ins across separate history folders rather than one extensive report containing everything at once.

When a file is checked out for editing in SCCS, it is locked so that anyone else cannot edit it. This prevents changes from being overwritten by other users and limits development since only one user can edit the file simultaneously.

The software supports branches that store sequences of changes to specific files - these can then be merged back into the original versions or with copies of other branched versions of the same parent branch.

Revision Control System (RCS)

The Revision Control System was released in 1982 as an alternative to the closed-source Source Code Control System. Developed by Walter Tichy and written in C, RCS was released under the GNU General Public License, making it suitable for use in open-source projects.

”RCS manages revisions of text documents, in particular source programs, documentation, and test data. It automates the storing, retrieval, logging and identification of revisions.” – Walter Tichy [11]

Architecture

RCS shares many traits with its predecessor[9], including:

- Handling revisions on a file-by-file basis.
- Changes across multiple files can’t be grouped together into an atomic commit.
- Tracked files are intended to be modified by one user at a time.
- No network functionality.
- Revisions for each tracked file are stored in a corresponding history file.
- Basic branching and merging of revisions within individual files.

Upon checking a file into RCS for the first time, a corresponding history file is created in the local `./RCS/` directory. The history file is postfixed with a `,v`, so a file named `test.txt` would be tracked by a file called `test.txt,v`[9].

RCS employs a reverse-delta scheme for storing file changes. When a file is checked in, the history file contains a complete snapshot of its contents. Subsequent modifications and check-ins result in RCS calculating a single delta—the difference between the new version of that specific revision and the previously recorded version - and saving it along with an older snapshot if necessary.

The reverse-delta scheme functions by checking out an earlier revision from the newest version and applying consecutive deltas until reaching the desired revision. Starting from the newest version allows quick checkout times, as the current revisions' snapshots are readily accessible.

However, when attempting to access older versions beyond just one recent update (e.g., 1 or 2 old versions), the process becomes considerably more complex. This is because these older versions' snapshots must be calculated against each other before they can be applied to the newest version.

Unlike RCS, SCCS maintains a consistent checkout time for any revision. Additionally, RCS history files do not store checksums, meaning file integrity cannot be guaranteed.

2.2.2 Centralized Version Control Systems (CVCS)

A Centralised Version Control System is a system that enables developers to work together on the same project by storing the primary copy of files in a central repository. This system keeps track of all files and saves information in the local repository. CVCS is called centralised because there is only one central server or repository.

The server maintains a complete record of issues, while clients only maintain a local copy of the shared documents. All developers make their modifications to the repository through checkout. However, only the last version of the files is retrieved from the server, meaning any modifications made will automatically be shared with other developers.

Users can modify in parallel with their local copy of shared documents and sync with the central server to release their contributions and make them visible to other collaborators. However, because centralised version control systems rely on one repository that includes the correct project version, they must restrict write access so that only trusted contributors can commit modifications.

CVCS has some challenges, such as if the central server is inaccessible, users cannot merge their work or save the released modifications. Also, if the central repository is corrupted, everything will be lost. Contributors must be the ones who have writing permissions to perform basic tasks, such as reverting modifications to a previous state, creating or merging branches, or releasing modifications with complete revision history. This limitation affects participation and authorship for new contributors.

So, the main drawbacks of using CVCS are that it requires a network connection to work on the source code, developers must order to contribute to a project, and a single point of failure is an issue when using one server.

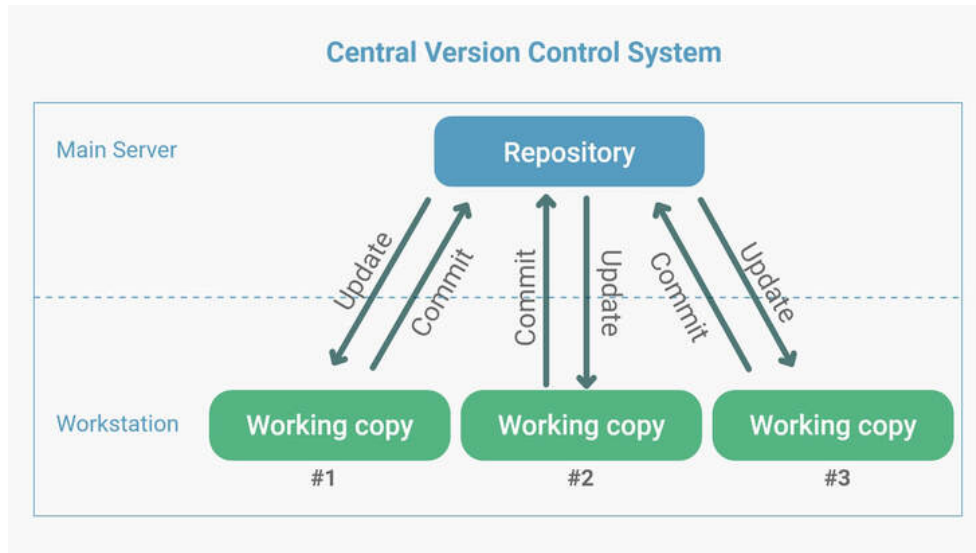


Figure 2.3: Centralized Version Control System (CVCS)

Concurrent Versions System (CVS)

Dick Grune developed the Concurrent Versions System in 1986 to introduce a networking element to version control. Written in C, CVS became the first widely used VCS tool that enabled multiple users to work on the same project simultaneously from different locations. This innovation began the second generation of VCS tools and facilitated collaboration among geographically dispersed development teams.

Architecture

CVS functions as a frontend for RCS, providing a set of commands to interact with files in a project while utilising the RCS history file format and commands behind the scenes. CVS allowed multiple developers to check out and work on duplicate files for the first time in history by employing a centralised repository model.

When a project is imported into CVS, each file is converted into a `.v` history file and stored in a central directory referred to as a `module`[9]. Generally, the repository resides on a remote server accessible via a local network or the Internet.

A developer checks out a copy of the module, which is copied to a working directory on their local machine. No files are locked during this process, allowing an unlimited

number of developers to check out the module simultaneously. In addition, developers can modify their checked-out files and commit changes as needed.

When a developer commits a change, other developers must update their working copies through a (usually) automated merge process before committing their changes. Occasionally, merge conflicts require manual resolution before a commit can be made. Therefore, CVS also offers the capability to create and merge branches.

Perforce Helix Core

Perforce Helix Core is a proprietary VCS developed, owned, and maintained by Perforce Software Inc., Written in `C` and `C++`; it was initially released in 1995. Although primarily designed for a centralised model, it also offers a distributed model option. Helix Core is commonly used by large companies managing substantial content with large binary files, such as in the video game development industry. Although typically cost-prohibitive for smaller projects, Perforce provides a free version for teams of up to five developers.

Architecture

Perforce Helix Core utilises a server/client model. The server, a process called `p4d`, listens for incoming client connections on a designated port, typically port `1666`. The client, a process called `p4`, is available in both command-line and GUI variants. Users run the `p4` client to connect to the server and issue commands. Support for various programming language APIs, including Python and Java, enables automated issuance and processing of Helix Core commands via scripts. Integrations are also available for IDEs like Eclipse and Visual Studio, allowing users to work with version control within those tools.

The Helix Core Server manages repositories called depots, which store files in directory trees similar to Apache Subversion (SVN). Clients can check out sets of files and directories from the depots into local copies called **workspaces**. The atomic unit used to group and track changes in Helix Core depots is called the **changelist**, which is analogous to Git commits. In addition, Helix Core implements two similar forms of branching: **branches** and **streams**. While branches represent separate lines of development history, a stream is a branch with added functionality that Helix Core uses to provide recommendations on best merging practices throughout the development process.

When a file is added for tracking, Helix Core classifies it using a `file type` label. The two most commonly used file types are text and binary. For binary files, the entire file content is stored each time the file is saved. This is a common VCS tactic

for handling binary files, which are not amenable to the normal merge process, as manual conflict resolution is typically impossible.

Only the deltas (changes between revisions) are stored for text files. Text file history and deltas are stored using the Revision Control System (RCS) format, which tracks each file in a corresponding `,v` file in the server depot. This is similar to Concurrent Versions System (CVS), which also leverages RCS file formats for preserving revision history. Files are often compressed using `gzip` when added to the depot and decompressed when synced back to the workspace.

Apache Subversion (SVN)

Subversion was created in 2000 by CollabNet Inc. and is now maintained by the Apache Software Foundation. Written in `C`, it was designed to offer a more robust centralised solution than Concurrent Versions System (CVS)[9].

Architecture

Similar to Concurrent Versions System (CVS), Subversion employs a centralised repository model, requiring remote users to have a working network connection to commit their changes to the central repository. However, in contrast to CVS, where a commit operation could fail midway due to a network outage and leave the repository in a corrupted and inconsistent state, Subversion ensures the repository remains consistent. Additionally, a Subversion commit or revision can include multiple files and directories, enabling users to track sets of related changes together as a grouped unit, as opposed to previous storage models that tracked changes separately for each file.

Subversion utilises a storage model called **FSFS (File System atop the File System)** for tracked files. The model's name originates from its database structure, which mirrors the operating system's file and directory structure. However, subversion's unique feature is its filesystem design, which tracks not only files and directories but also different versions of these files and directories as they change over time. This creates a filesystem with an added time dimension. Moreover, Subversion treats folders as first-class citizens, allowing empty folders to be committed, unlike other systems such as Git, where empty folders go unnoticed.

When a Subversion repository is created, a (nearly) empty database of files and folders is established [9]. Next, a directory called **db/revs** is created to store all revision tracking information for the committed files. Each commit, which may include changes to multiple files, is stored in a new file in the **revs** directory, named with a sequential numeric identifier starting with 1. When a file is committed for the first time, its entire content is stored. Subsequent commits of the same file store only

the changes, also known as `diffs` or `deltas`, to save space. Additionally, `deltas` are compressed using `lz4` or `zlib` compression algorithms to reduce their size further.

By default, this approach is only employed to a certain extent. Although storing file deltas instead of the entire file each time saves storage space, it adds time to checkout and commit operations, as all the deltas must be combined to recreate the file's current state. For this reason, Subversion stores up to 1023 deltas per file before saving a new complete copy of the file, striking a balance between storage and speed.

2.2.3 Distributed Version Control Systems (DVCS)

Distributed Version Control Systems (DVCS) were developed to address the limitations of Centralized Version Control Systems (CVCS), enabling more effective branching and merging, seamless local VCS operations, and improved collaboration among developers. Due to these limitations associated with Centralized Version Control Systems (CVCS), Distributed Version Control Systems have become widely adopted in Open Source Software (OSS) projects.

DVCS is designed to function in two ways: it maintains file histories locally on each device and can synchronise local user modifications with servers when necessary, enabling the sharing of these modifications with others. In DVCS, developers can work independently or collaboratively on the same project, as they have access to all necessary repositories. In addition, any repository can be cloned from another, meaning no repository holds greater importance than others.

In response to the need for more advanced versioning of software artefacts, several Distributed Version Control Systems, such as Mercurial, and Git, emerged in the software field. Many OSS projects have widely adopted these tools. DVCS operations are significantly faster than those in CVCS, as they are performed locally, while CVCS operations require remote connections. Some experts believe that distributed systems will eventually replace centralised ones due to their suitability for more extensive projects involving independent developers seeking full functionality, even without a network connection. In addition, distributed systems offer advantages such as saving earlier drafts of work without publicly releasing or sharing them with others.

Collaboration between team members and the ability for individual developers to serve as servers or clients are essential features of version control systems, allowing developers to work on source code without being connected to a central or remote repository.

However, DVCS introduces some challenges, including the lack of a coherent version numbering system due to the absence of a centralised versioning server. Instead, DVCS relies on hash modifications or a unique GUID. This lack of a central server also complicates system backup.

The two most common criticisms of DVCS are the unavailability of pessimistic locks and weak support for binary files. Despite these disadvantages, the transition from centralised to decentralised version control systems continues, as developers can work offline and efficiently work incrementally. This flexibility enables developers to assume multiple roles, such as developing new tasks or fixing errors, leading to exploratory coding and greater freedom in their workflow while maintaining control over their code's release schedule.




Figure 2.4: Distributed Version Control System (DVCS)

Git

Git, created in 2005 by Linus Torvalds (also the creator of Linux), is primarily written in C with some shell scripts. Git was initially developed for the Linux codebase and has since become the most popular VCS in use today due to its features, flexibility, and speed. Torvalds explains that Git is a robust set of tools with many options, and its usage is often determined by what works best for collaboration rather than technical limitations.

"You can do a lot of things with Git, and many of the rules of what you should do are not so much technical limitations but are about what works well when working together with other people. So Git is a very powerful set of tools, and that can not only be overwhelming at first, it also means that you can often do the same (or similar) things different ways, and they all 'work.'" – Linus Torvalds [1]

Git repositories are commonly hosted on local servers and cloud services, forming the backbone of a broad set of DevOps tools available from popular service providers, including GitHub, BitBucket, GitLab, and many others [9].

Architecture

As a Distributed Version Control Systems (DVCS), Git ensures that no repository copy needs to be designated as the centralised copy—instead, all copies are created equal. This contrasts with second-generation VCS, which relies on a centralised copy for users to check in and out.

This design allows developers and coding partners to share changes directly with each other before merging their changes into an official branch, fostering a flexible distributed workflow for team collaboration.

Moreover, developers can commit changes to their local copy of the repository without other repositories knowing about it. This enables commits without a network or internet connection, allowing developers to work offline until they are ready to push their changes to other repositories for review, testing, or deployment.

When a file is added for tracking with Git, it is compressed using the `zlib` compression algorithm and hashed using a `SHA-1` hash function [9]. This generates a unique hash value corresponding to the file content, which Git stores in an object database located in the hidden `.git/objects` folder. These files, called **Git blobs**, are created each time a new file (or a changed version of an existing file) is added to the repository.

Git uses a staging index that functions as a temporary space for changes being readied for a commit. When changes are set to be committed, their compressed data is referenced in a unique index file, appearing as a tree object. **Trees** in Git link blob objects to actual file names, file permissions, and connections to other trees, signifying the status of a specific collection of files and directories. After all associated changes have been staged for commit, the index tree can be committed to the repository, generating a commit object within the Git object database [9].

A commit denotes the primary tree for a specific revision, and the commit author, email address, date, and a descriptive commit message. Each commit also retains a reference to its preceding commit/commits, constructing a record of the project's evolution.

Git objects, including blobs, trees, and commits, are all compressed, hashed, and saved in the object database based on their respective hash values. These standalone objects avoid using diffs for space conservation, making Git highly efficient since the complete content of every file revision is readily available as an individual object.

Nonetheless, particular operations, such as pushing commits to a remote repository, storing an excessive number of objects, or manually executing Git's garbage collection command, may prompt Git to reorganise objects into pack files. This packing procedure compresses inverse diffs to remove duplicate content and minimise size. This leads to the creation of **.pack** files containing the object data, each paired with a corresponding **.idx** (index) file that references the packed objects and their positions within the pack file. These pack files are transmitted across the network when branches are pushed to or fetched from remote repositories. When pulling or retrieving branches, the pack files are decompressed to generate loose objects in the object repository.

Mercurial

Mercurial, created in 2005 by Matt Mackall and written in **Python**, initially aimed to host the codebase for Linux, but Git was chosen instead [9]. As the second most popular distributed VCS after Git, Mercurial is used far less frequently.

Architecture

Comparable to Git, Mercurial is a Distributed Version Control Systems (DVCS) that enables multiple developers to work on separate copies of the same project. Although Mercurial utilises many similar technologies as Git, such as compression and **SHA-1** hashing, it does so in distinct ways.

When a new file is tracked in Mercurial, a corresponding **revlog** (revision log) file is generated for it in the hidden **.hg/store/data/** directory. **Revlog** files can be viewed as advanced versions of the history files employed by older VCSs such as Concurrent Versions System (CVS), Revision Control System (RCS), and Source Code Control System (SCCS).

Unlike Git, which generates a new blob for each version of every staged file, Mercurial merely creates a new entry in the revlog for that file. To conserve space, each new entry contains only the delta (modifications) from the preceding version. Once a certain number of deltas is achieved, a complete snapshot of the file is stored again, minimising lookup time when applying numerous deltas to reconstruct a specific file revision.

These file revlogs are named to correspond with the files they monitor but are postfixed with **.i** and **.d** extensions. The **.d** files hold the compressed **delta** content, while the **.i** files function as **indexes** to locate distinct revisions within the **.d** files swiftly. For small files with few revisions, both the **indexes** and **content** are stored in **.i** files. Revlog file entries are compressed for efficiency and hashed for identification, with the hash values referred to as **nodeids**.

Chapter 3

Design

This chapter examines various data structures and algorithms with key aspects that may make them suitable for incorporation into a **Version Control System**. We also assess what metrics are relevant to facilitate the comparison of these data structures and algorithms.

3.1 Data Structures

Data structures are objects that can be used to store, organize, and manipulate large amounts of data. They play a crucial role in computer science and are fundamental building blocks of many software systems, including **Version Control Systems**. The choice of data structure is critical to the overall performance and scalability of a **Version Control System**.

In order to evaluate the suitability of a data structure for implementation in a **Version Control System**, it is necessary to consider several core aspects. Firstly, **operational efficiency**, which refers to the time and space complexity of the basic operations performed by the data structure, is a crucial consideration. The efficiency of these operations can have a significant impact on the overall performance of the **Version Control System**.

Another important aspect is the data structure's **structural specificity**, which encompasses how data is stored and organized within the structure. This is critical because the structural specifics can affect the ease of implementation and the ability to efficiently perform operations such as **insertions**, **deletions**, and **updates**.

Lastly, the **implementation details** must be considered, including the ease of implementation and compatibility with the programming language. A data structure that is straightforward to implement and easy to maintain/iterate upon will result in a more streamlined and efficient **Version Control System**.

3.1.1 Linked List

A **Linked List** is a linear collection of data elements, called nodes, each pointing to the next node by means of a pointer. The **Linked List** is the most sought-after data structure when it comes to handling dynamic data elements [5].

There are two main types of **Linked Lists**: Singly-linked lists (SLL) and Doubly-linked lists (DLL), but we will only be considering the Doubly-linked lists (DLL) when we reach the Implementation chapter.

Singly-linked lists (SLL)

- SLL nodes contain two fields: **data** field and **next** pointer field.
- Traversal of a SLL can be done using the **next** pointer field only. Meaning, the SLL can be traversed in only one direction, from the first node to the last node.
- The SLL occupies less memory than a DLL because it does not contain a **prev** pointer field.
- SLL is preferred over DLL when it comes to the execution of stack and queue operations.
- SLL is also preferred over DLL to save memory when a searching operation is not required.

Table 3.1: Efficiency Analysis of Singly-Linked List Operations

Operation	Worst Case	Average Case
Access	$O(n)$	$O(n)$
Search	$O(n)$	$O(n)$
Insert (at Head)	$O(1)$	$O(1)$
Delete (at Head)	$O(1)$	$O(1)$
Insert (at Current)	$O(1)$	$O(1)$
Delete (at Current)	$O(1)$	$O(1)$
Insert (at Tail)	$O(n)$	$O(n)$
Delete (at Tail)	$O(n)$	$O(n)$




Figure 3.1: Singly Linked List (SLL) [12]

Doubly-linked lists (DLL)

- DLL nodes contain three fields: **data** field, **prev** pointer field and **next** pointer field.
- Traversal of a DLL can be done using the **next** pointer field or the **prev** pointer field. Meaning, the DLL can be traversed in both directions, from the first node to the last node and vice versa.
- The DLL occupies more memory than a SLL because it contains a **prev** pointer field.

Table 3.2: Efficiency Analysis of Doubly-Linked List Operations

Operation	Worst Case	Average Case
Access	$O(n)$	$O(n)$
Search	$O(n)$	$O(n)$
Insert (at Head)	$O(1)$	$O(1)$
Delete (at Head)	$O(1)$	$O(1)$
Insert (at Current)	$O(1)$	$O(1)$
Delete (at Current)	$O(1)$	$O(1)$
Insert (at Tail)	$O(1)$	$O(1)$
Delete (at Tail)	$O(1)$	$O(1)$




Figure 3.2: Doubly Linked List (DLL) [12]

Advantages

- Efficient insertion and deletion operations at the beginning, end, and middle of the list.
- Easy traversal in both directions, allowing for simpler and more flexible algorithms.
- Dynamic size, allowing the list to grow and shrink as needed during runtime.
- No need for contiguous memory allocation, which allows for better memory utilisation.
- Simplified implementation compared to more complex data structures.

Disadvantages

- Higher memory overhead due to the storage of two pointers for each node.
- Slower random access compared to arrays or hash tables, as elements must be traversed sequentially.
- No inherent support for efficient searching, leading to linear search times.

Summary

A **Doubly Linked List** could be used as the core data structure, but it has some limitations. The sequential nature of the data structure makes it easy to maintain a linear history of changes and revert to previous versions. However, the lack of efficient searching and random access capabilities can slow down operations when dealing with large repositories or complex branching scenarios.

3.1.2 Binary Tree

A **Binary Tree** is a hierarchical data structure in which each node has at most two child nodes, arranged in a way that the value of the node to the left is less than or equal to the parent node and the value of the node to the right is greater than or equal to the parent node. This ordering property ensures efficient search, insertion, and deletion operations. There are several types of binary trees, such as **binary search trees**, **AVL trees**, and **red-black trees**, each with different balancing mechanisms to maintain tree height and performance.

Each node in a **Binary Tree** contains the following elements:

1. Data: The data stored in the node.
2. Left child: A pointer to the left child node.
3. Right child: A pointer to the right child node.

Table 3.3: Efficiency Analysis of Binary Tree Operations

Operation	Worst Case	Average Case
Search	$O(n)$	$O(\log n)$
Insert	$O(n)$	$O(\log n)$
Delete	$O(n)$	$O(\log n)$

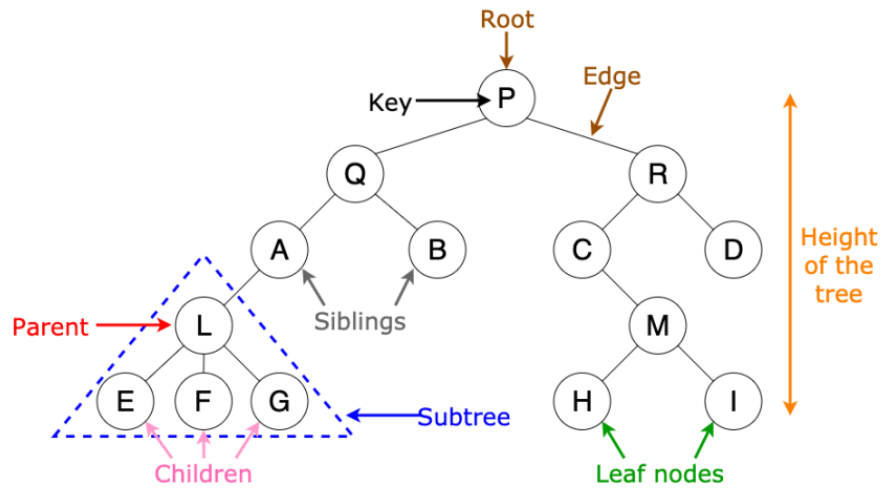


Figure 3.3: Binary Tree [2]

Advantages

- Efficient search, insertion, and deletion operations when the tree is balanced.
- Hierarchical structure allows for natural representation of hierarchical relationships or data with partial order.
- Can be easily traversed in various orders (e.g., inorder, preorder, and postorder) to suit different needs.
- No need for contiguous memory allocation, which allows for better memory utilisation.
- Provides the foundation for more advanced tree structures, like B-trees or trie, that can be used for advanced indexing or searching.

Disadvantages

- Unbalanced trees can lead to degraded performance.
- Requires more memory overhead compared to linear data structures due to the storage of pointers for each node.
- Less suitable for representing non-hierarchical or unordered data.
- Can be more complex to implement and maintain compared to simpler data structures.

Summary

A **Binary Tree** could be used as the core data structure, but it may not be the most suitable choice due to its hierarchical nature. While it can efficiently store and manage version history when dealing with linear or partially ordered data, version control systems often require support for complex branching and merging scenarios, which may not be well-suited for a binary tree.

Additionally, the need for balancing mechanisms to maintain tree height and performance can add complexity to the implementation and maintenance of the system. **Directed Acyclic Graphs** or other more advanced data structures might be more appropriate for handling complex functionality.

3.1.3 Hash Table

A **Hash Table** is a data structure that uses a hash function to map keys to their corresponding values. It is an efficient way to implement an associative array, where keys are used to look up values.

The basic idea behind a **Hash Table** is to use a hash function to map a key to an index in an array, called a bucket, where the corresponding value can be found or stored. The process of mapping a key to an index is called **hashing**.

Each element in a **Hash Table** consists of:

1. **Key**: This is the value used to look up a corresponding element in the **Hash Table**.
2. **Value**: This is the value associated with the key that is stored in the **Hash Table**.

When a new element is added to a **Hash Table**, the key is passed through a **hash** function which produces an **index** (also called a hash value or bucket) where the element is stored. When a value is to be retrieved, the key is passed through the same hash function, and the resulting **index** is used to look up the corresponding value in the **Hash Table**.

Table 3.4: Efficiency Analysis of Hash Table Operations

Operation	Worst Case	Average Case
Search	$O(n)$	$\Theta(1)$
Insert	$O(n)$	$\Theta(1)$
Delete	$O(n)$	$\Theta(1)$




Figure 3.4: Hash Table [8]

Advantages

- Fast average-case performance for search, insertion, and deletion operations.
- Supports efficient key-based lookups and direct access to values.
- Can be easily resized to accommodate a growing number of key-value pairs, maintaining constant-time complexity.
- Suitable for storing unordered or non-hierarchical data.
- Provides a foundation for more advanced data structures, like distributed hash tables or bloom filters, used in various applications.

Disadvantages

- Requires a good hash function to ensure uniform key distribution and avoid performance degradation due to collisions.
- Higher memory overhead compared to linear data structures, as the underlying array needs to be larger than the number of stored key-value pairs to maintain performance.
- No inherent support for ordered traversal or range queries, as keys are not stored in a sorted manner.

Summary

While **Hash Tables** can provide fast and efficient key-based lookups, they may not be the most suitable data structure for the core of a **Version Control System**. The lack of inherent support for ordered traversal or range queries can make it difficult to efficiently handle complex branching and merging scenarios that are common in version control systems.

Directed Acyclic Graphs or other more advanced data structures are often better suited for handling complex branching and merging operations, as they can provide more efficient support for ordered traversal and range queries. Additionally, Version Control Systems typically need to maintain relationships between revisions, which is not a natural fit for the unordered nature of hash tables.

In summary, although **Hash Tables** can provide fast key-based lookups and efficient performance in certain scenarios, they may not be the most suitable or scalable choice for the core data structure of a version control system due to their unordered nature and lack of support for ordered traversal or range queries.

3.1.4 Directed Acyclic Graph (DAG)

A **Directed Acyclic Graph (DAG)** is a data structure that consists of nodes connected by directed edges, forming a graph with no cycles. This means that it is impossible to start at a node and follow a sequence of directed edges that leads back to the same node. DAGs are particularly useful for representing dependencies, partial orders, or processes that have a specific sequence of events. For example, in the context of a version control system, a DAG can be used to represent the history of changes made to a project, with nodes representing commits and edges representing parent-child relationships between commits.




Figure 3.5: Directed Acyclic Graph (DAG) [10]

Advantages

- Efficiently represents complex branching and merging scenarios common in Version Control Systems.
- Can naturally model dependencies, partial orders, or processes with a specific sequence of events.
- Enables efficient algorithms for topological sorting, which can be useful for ordering commits or resolving dependencies.
- No need for contiguous memory allocation, which allows for better memory utilisation.
- Provides a more expressive and flexible data structure compared to linear or hierarchical structures, making it suitable for various applications.

Disadvantages

- Can be more complex to implement and maintain compared to simpler data structures.
- Requires more memory overhead compared to linear data structures due to the storage of multiple pointers for each node.
- Finding the shortest or most efficient path between nodes can be computationally expensive in large graphs.
- No inherent support for fast key-based lookups or direct access to values, as nodes are not indexed by a specific key.

Summary

Directed Acyclic Graphs are particularly well-suited for use as the core data structure in a version control system. Their ability to efficiently represent complex branching and merging scenarios allows for more expressive and flexible management of project history. Additionally, the natural modelling of dependencies and partial orders enables efficient algorithms for tasks such as topological sorting and dependency resolution, which are common in version control systems. Overall, the advantages of DAGs in representing complex relationships and dependencies make them a suitable and scalable choice for the core data structure of a version control system.

3.2 Algorithms

3.2.1 Searching

Search algorithms are computational methods to locate specific items, solve problems, or explore and analyse data structures. They are essential tools in computer science, artificial intelligence, and various other domains, as they facilitate efficient processing and management of large datasets. By systematically searching through data or traversing complex structures, these algorithms identify target elements, optimal solutions, or specific patterns.

Linear Search

Linear Search is a simple and straightforward search algorithm that works on unsorted and sorted lists. It sequentially checks each element in the list until the desired element is found or the end of the list is reached [7]. The algorithm compares each element in the list with the target value, starting from the first element and moving through the list one element at a time. Linear Search has an average case time complexity of $O(n)$ and a worst case time complexity of $O(n)$, where n is the number of elements in the list.




Figure 3.6: Linear Search Algorithm [6]

Advantages

- Simple and easy to implement.
- Works on unsorted and sorted lists.
- No additional memory requirements or data structure modifications needed.
- Performs well for small data sets or when the target element is near the beginning of the list.

- Can be used as a building block for more advanced search algorithms.

Disadvantages

- Inefficient for large data sets, as it has to check each element in the list.
- Slower than other search algorithms for sorted lists, such as binary search.
- Does not take advantage of any existing order in the list.
- May not be suitable for real-time or performance-critical applications.

Binary Search

Binary Search is an efficient search algorithm that works on sorted lists [4]. It repeatedly divides the list in half, comparing the middle element with the target value. If the middle element is equal to the target, the search is successful. If the target value is less than the middle element, the search continues in the left half of the list, otherwise in the right half. This process is repeated until the target value is found or the remaining list becomes empty. Binary Search has an average case time complexity of $O(\log n)$ and a worst case time complexity of $O(\log n)$, where n is the number of elements in the list.




Figure 3.7: Binary Search Algorithm [3]

Advantages

- Efficient search algorithm for sorted lists, with a time complexity of $O(\log n)$.
- Takes advantage of the existing order in the list.

- Requires fewer comparisons than linear search for large data sets.
- Can be implemented iteratively or recursively.
- Provides a foundation for more advanced search algorithms, like interpolation search.

Disadvantages

- Requires the list to be sorted in ascending order beforehand.
- Not suitable for unsorted lists or lists with frequently changing data.
- Can be more complex to implement compared to linear search.
- Does not work well with large lists stored on slow access media (e.g., hard drives) due to multiple random access operations.

Depth First Search (DFS)

Depth First Search is a graph traversal algorithm that explores as far as possible along a branch before backtracking. It can be implemented using recursion or an explicit stack data structure. Starting from a source node, DFS visits a node and marks it as visited. Then, it recursively explores each unvisited neighbour of the current node, treating the neighbor as the new current node. The algorithm continues this process until all reachable nodes have been visited. DFS has an average case time complexity of $O(V + E)$ and a worst case time complexity of $O(V + E)$, where V is the number of vertices and E is the number of edges in the graph.

Advantages

- Can be used to find connected components, cycles, and paths in a graph.
- Effective for searching large, sparsely connected graphs.
- Can be implemented using recursion or an explicit stack data structure.
- Visits nodes in a linear and natural order.
- Can be adapted for other graph traversal problems, like topological sorting.




Figure 3.8: Depth First Search Algorithm [13]

Disadvantages

- May consume a large amount of memory for deep graphs when implemented recursively.
- Can get stuck in cycles if not properly implemented with visited node tracking.
- Does not always find the shortest path in weighted graphs.
- May be less intuitive than Breadth First Search for certain problems.

Breadth First Search (BFS)

Breadth First Search is a graph traversal algorithm that visits all nodes at the same level before moving on to the next level. It can be implemented using a queue data structure. Starting from a source node, BFS visits and marks the node as visited. Then, it enqueues all unvisited neighbors of the current node. Finally, the algorithm dequeues the next node from the front of the queue and repeats the process until the queue is empty, ensuring that all reachable nodes have been visited. BFS has an average case time complexity of $O(V + E)$ and a worst case time complexity of $O(V + E)$, where V is the number of vertices and E is the number of edges in the graph.

Advantages




Figure 3.9: Breadth First Search Algorithm [13]

- Can be used to find the shortest path in unweighted graphs or determine the level of each node from the source node.
- Effective for searching large, densely connected graphs.
- Can be implemented using a queue data structure, avoiding recursion.
- Visits nodes in a level-wise order, which can be more intuitive for certain problems.
- Can be adapted for other graph traversal problems, like bipartite graph checking.

Disadvantages

- Can consume a large amount of memory for densely connected graphs due to the use of a queue.
- Not well-suited for finding all paths or connected components in a graph.
- Does not work efficiently for searching large, sparsely connected graphs.
- Can be slower than **Depth First Search** for certain problems, such as finding cycles or connected components.

3.2.2 Diffing Algorithms

Diffing algorithms, also known as difference algorithms or delta algorithms, are specialised computational methods designed to identify and highlight the differences between two sets of data. These algorithms compare input sequences, such as text, code, or binary data, to efficiently determine the changes needed to transform one sequence into the other. By detecting additions, deletions, and modifications, diffing algorithms enable the tracking and management of revisions in a concise and comprehensible manner.

Myers Diff Algorithm

The **Myers diffing algorithm**, developed by Eugene W. Myers, is an efficient algorithm for comparing two sequences (such as lines in text files) and finding the shortest edit script that transforms one sequence into the other. The algorithm is based on the concept of an edit graph, where each node represents an element in the sequences, and the edges represent insertions, deletions, or matches. The algorithm finds the shortest path through this graph, also known as the **Longest Common Subsequence (LCS)** path, which represents the minimal set of edits required to transform one sequence into the other.

The time and space complexity of the **Myers Diff** algorithm is $O(ND)$, where N is the sum of the lengths of both inputs, and D is the size of the minimum edit script that converts one input to the other. When the number of differences is small, various optimisations can be applied to improve complexity up to $O(N \log N + D^2)$ time and $O(N)$ space.

Advantages

- Efficiently finds the shortest edit script, reducing the number of changes needed to transform one sequence into the other.
- Can handle large input sequences, as it has a linear space complexity.
- Can be adapted for various applications, such as comparing text files, source code, or DNA sequences.
- The algorithm is well-established and widely used, making it a reliable choice for many diffing tasks.

Disadvantages

- Can produce suboptimal diffs for specific cases, such as when comparing sequences with many common elements but in a different order.
- The algorithm's implementation can be complex and difficult to understand for those unfamiliar with graph-based algorithms.
- May not be the most efficient algorithm for all diffing tasks, as other algorithms may produce better results in specific cases.
- The quality of the output can be sensitive to the choice of the underlying **Longest Common Subsequence** algorithm.

Patience Diff Algorithm

The **Patience Diffing** algorithm, developed by Bram Cohen, is a diffing algorithm that aims to produce more human-readable diffs by focusing on finding the **Longest Increasing Subsequence** (LIS) in the two sequences. The algorithm first identifies unique matching lines (or elements) between the sequences and sorts them by their position in the first sequence. Then, it finds the LIS in the sorted list of positions from the second sequence. The algorithm recursively applies this process to the unmatched portions of the sequences until the entire diff is produced.

The time complexity of the **Patience Diff** algorithm is $O(N \log N)$, where N is the sum of the lengths of both inputs.

Advantages

- Produces more human-readable diffs, as it focuses on preserving the order of common elements.
- Can handle large input sequences, as it has a linearithmic time complexity.
- Can be used as a standalone algorithm or in combination with other diffing algorithms, such as the Myers algorithm, to improve the quality of the output.
- The algorithm is relatively simple and straightforward to implement compared to some other diffing algorithms.
- Works well for cases where the sequences have many common elements but in a different order.

Disadvantages

- May produce suboptimal diffs for specific cases, such as when comparing sequences with many non-unique elements.
- The algorithm's performance depends on the choice of the underlying LIS algorithm.
- May not be the most efficient algorithm for all diffing tasks, as other algorithms may produce better results in specific cases.
- Not as widely used or well-established as some other diffing algorithms, such as the `Myers algorithm`.

Chapter 4

Implementation

4.1 Data Structures

4.1.1 Revision Type

Before implementing any data structures, I decided to create a foundational component that would serve as a building block for the subsequent data structures and algorithms. This foundational component is the **Revision** interface and its implementation, which is designed to represent a single revision of a file in a Version Control System (VCS). The primary goal of this approach was to abstract the representation of a file revision, allowing for a consistent and modular way to manage revisions across the different data structures and algorithms being implemented and compared.

```
type Revision interface {  
    ID() int  
    Data() []byte  
}
```

The **Revision** interface defines two methods: one for obtaining the revision ID as an integer and another for accessing the revision data as a byte slice. The concrete implementation of this interface encapsulates the revision ID and data in two fields. This design enables us to consistently create and manipulate file revisions, regardless of the specific data structures or algorithms employed.

```
type standardRevision struct {  
    id    int  
    data []byte  
}
```

```

func (rev *standardRevision) ID() int {
    return rev.id
}

func (rev *standardRevision) Data() []byte {
    return rev.data
}

```

Three functions have been provided to create new revisions with varying initial data. The first function allows users to create a revision with a specified ID and data, while the second function creates a revision with the given ID and no data. Finally, the third function generates a revision with random data of a specified size.

```

func NewRevision(revisionID int, revisionData []byte) Revision {
    return &standardRevision{revisionID, revisionData}
}

func NewBlankRevision(revisionID int) Revision {
    return NewRevision(revisionID, nil)
}

func NewRandomRevision(revisionID int, dataSize int) Revision {
    data, err := generateRandomBytes(dataSize)
    if err != nil {
        log.Fatalf("error generating random bytes: %v", err)
    }
    return NewRevision(revisionID, data)
}

```

This `Revision` abstraction can be employed to simulate different scenarios or workloads for the data structures and algorithms under evaluation. For example, the function that generates a revision with random data of a specified size could be used to create many random revisions, allowing for the evaluation of each algorithm's performance in handling diverse and unpredictable data.

Tests

TestNewRevision This test case checks if the `NewRevision` function creates a new revision with the correct ID and data. It compares the ID and data of the created revision with the expected values and reports an error if they don't match.

TestNewBlankRevision This test case checks if the `NewBlankRevision` function creates a new revision with the correct ID and no data. It compares the ID of the created revision with the expected value and reports an error if it doesn't match. It also checks if the revision data is `nil`.

TestNewRandomRevision This test case checks if the `NewRandomRevision` function creates a new revision with the correct ID and random data of the specified size. It compares the ID and length of the data of the created revision with the expected values and reports an error if they don't match.

Benchmarks

BenchmarkNewRevision This benchmark measures the performance of the `NewRevision` function. It creates a new revision with the given ID and data repeatedly in a loop, and the testing framework records the time taken.

BenchmarkNewBlankRevision This benchmark measures the performance of the `NewBlankRevision` function. It creates a new revision with the given ID and no data repeatedly in a loop, and the testing framework records the time taken.

BenchmarkNewRandomRevision This benchmark measures the performance of the `NewRandomRevision` function. It creates a new revision with the given ID and random data of the specified size repeatedly in a loop, and the testing framework records the time taken.

BenchmarkGenerateRandomBytes This benchmark measures the performance of the `generateRandomBytes` function. It generates a slice of random bytes of the specified length repeatedly in a loop, and the testing framework records the time taken. If there's an error generating the random bytes, the benchmark stops with a fatal error.

4.1.2 Doubly Linked List

After establishing a solid foundation with the `Revision` interface, the next step in the implementation stage was to create a doubly linked list data structure that could store revisions. The primary reason for choosing a doubly linked list as one of the data structures for comparison was its flexibility and efficiency in insertion and deletion operations. Additionally, its dynamic nature allows for efficient memory utilisation.

The doubly linked list implementation consists of two primary components: the ‘DLLNode’ struct, which represents a node in the list, and the ‘DoublyLinkedList’ struct, which represents the list itself. The ‘DLLNode’ struct contains a Revision object along with pointers to the next and previous nodes in the list. On the other hand, the ‘DoublyLinkedList’ struct maintains pointers to the list’s head (first) and tail (last) nodes.

```
type DLLNode struct {
    Revision types.Revision
    Next     *DLLNode
    Prev     *DLLNode
}

type DoublyLinkedList struct {
    Head *DLLNode
    Tail *DLLNode
}
```

A set of methods were developed to facilitate various operations on the doubly linked list, such as adding and removing nodes, inserting nodes at specific positions, and checking for the presence of a node with a specified revision ID. These methods enable us to easily manipulate the doubly linked list and provide a way to evaluate its performance characteristics under different scenarios.

```
// AssignHead sets the given node as the head of the doubly linked list.
func (dll *DoublyLinkedList) AssignHead(node *DLLNode)

// AssignTail sets the given node as the tail of the doubly linked list.
func (dll *DoublyLinkedList) AssignTail(node *DLLNode)

// InsertPrior inserts a new node before the specified node in the doubly
// linked list.
func (dll *DoublyLinkedList) InsertPrior(node, nodeToInsert *DLLNode)

// InsertSubsequent inserts a new node after the specified node in the
// doubly linked list.
func (dll *DoublyLinkedList) InsertSubsequent(node, nodeToInsert *DLLNode)

// InsertAtPosition inserts a new node at the specified position in the
// doubly linked list.
func (dll *DoublyLinkedList) InsertAtPosition(position int, nodeToInsert
    *DLLNode)
```

```

// RemoveNodesWithID removes all nodes with the specified revision ID from
// the doubly linked list.
func (dll *DoublyLinkedList) RemoveNodesWithID(id int)

// Remove removes the specified node from the doubly linked list.
func (dll *DoublyLinkedList) Remove(node *DLLNode)

// ContainsNodeWithID returns true if the doubly linked list contains a
// node with the specified revision ID.
func (dll *DoublyLinkedList) ContainsNodeWithID(id int) bool

```

Tests

TestNewDoublyLinkedList This test case checks if the `NewDoublyLinkedList` function creates a new doubly linked list with a nil Head and Tail. It asserts that the Head and Tail of the created list are nil and reports an error if they're not.

TestDoublyLinkedList_AssignHead This test case checks if the `AssignHead` method assigns the given node as the head of the doubly linked list. It assigns the node as the head of the list and asserts that the list's Head and Tail are both equal to the node. It reports an error if they're not.

TestDoublyLinkedList_AssignTail This test case checks if the `AssignTail` method assigns the given node as the tail of the doubly linked list. It assigns the node as the tail of the list and asserts that the list's Head and Tail are both equal to the node. It reports an error if they're not.

TestDoublyLinkedList_InsertPrior This test case checks if the `InsertPrior` method inserts a new node before the specified node in the doubly linked list. It inserts a new node before the specified node and asserts that the list's Head, Tail, and node relationships are correct. It reports an error if they're not.

TestDoublyLinkedList_InsertSubsequent This test case checks if the `InsertSubsequent` method inserts a new node after the specified node in the doubly linked list. It inserts a new node after the specified node and asserts that the list's Head, Tail, and node relationships are correct. It reports an error if they're not.

TestDoublyLinkedList_InsertAtPosition This test case checks if the `InsertAtPosition` method inserts a new node at the specified position in the doubly linked list. It inserts a new node at the specified position and asserts that the list's Head, Tail, and node relationships are correct. It reports an error if they're not.

TestDoublyLinkedList_RemoveNodesWithID This test case checks if the `RemoveNodesWithID` method removes all nodes with the specified revision ID from the doubly linked list. It removes all nodes with the specified revision ID and asserts that the list's Head, Tail, and node relationships are correct. It reports an error if they're not.

TestDoublyLinkedList_ContainsNodeWithID This test case checks if the `ContainsNodeWithID` method returns true if the doubly linked list contains a node with the specified revision ID. It checks if the list contains nodes with specific revision IDs and asserts that the result is correct. It reports an error if the result is not as expected.

TestDoublyLinkedList_Remove This test case checks if the `Remove` method removes the specified node from the doubly linked list. It removes the specified node and asserts that the list's Head, Tail, and node relationships are correct. It reports an error if they're not.

Benchmarks

BenchmarkDoublyLinkedList_InsertPrior This benchmark case checks the performance of the `InsertPrior` method. It inserts a new node before the specified node in the doubly linked list and reports the time taken to perform the operation. The benchmark is run for a number of iterations and the average time taken is reported.

BenchmarkDoublyLinkedList_InsertSubsequent This benchmark case checks the performance of the `InsertSubsequent` method. It inserts a new node after the specified node in the doubly linked list and reports the time taken to perform the operation. The benchmark is run for a number of iterations and the average time taken is reported.

BenchmarkDoublyLinkedList_InsertAtPosition This benchmark case checks the performance of the `InsertAtPosition` method. It inserts a new node at the specified position in the doubly linked list and reports the time taken to perform the operation. The benchmark is run for a number of iterations and the average time taken is reported.

BenchmarkDoublyLinkedList_RemoveNodesWithID This benchmark case checks the performance of the `RemoveNodesWithID` method. It removes all nodes with the specified revision ID from the doubly linked list and reports the time taken to perform the operation. The benchmark is run for a number of iterations and the average time taken is reported.

BenchmarkDoublyLinkedList_ContainsNodeWithID This benchmark case checks the performance of the `ContainsNodeWithID` method. It checks if the list contains nodes with specific revision IDs and reports the time taken to perform the operation. The benchmark is run for a number of iterations and the average time taken is reported.

BenchmarkDoublyLinkedList_RevisionConstruction This benchmark case measures the performance of constructing doubly linked lists with varying numbers of revisions and different data sizes for each revision. This benchmark helps to understand the efficiency of the code when dealing with different scenarios.

```
const (
    smallNumRevisions = 200
    mediumNumRevisions = 2000
    largeNumRevisions = 20000
    largeDataSize      = 10000 // 10 Kb
    mediumDataSize      = 1000  // 1 Kb
    smallDataSize       = 100   // 100 bytes
)

func BenchmarkDoublyLinkedList_RevisionConstruction(b *testing.B) {
    for _, numRevisions := range []int{smallNumRevisions,
        mediumNumRevisions, largeNumRevisions} {
        for _, dataSize := range []int{smallDataSize, mediumDataSize,
            largeDataSize} {
            b.Run(fmt.Sprintf("numRevisions=%d,dataSize=%d", numRevisions,
                dataSize), func(b *testing.B) {
                for i := 0; i < b.N; i++ {
                    list := NewDoublyLinkedList()
                    for j := 0; j < numRevisions; j++ {
                        revision := types.NewRandomRevision(j, dataSize)
                        list.InsertAtPosition(j+1, &DLLNode{Revision: revision})
                    }
                }
            })
        }
    }
}
```



```

    }
  }
}

```

1. The benchmark iterates over two slices: one containing the number of revisions (`smallNumRevisions`, `mediumNumRevisions`, and `largeNumRevisions`) and the other containing the data size for each revision (`smallDataSize`, `mediumDataSize`, and `largeDataSize`).
2. For each combination of the number of revisions and data size, it runs a sub-benchmark using `b.Run`. The `fmt.Sprintf` function is used to create a descriptive name for the sub-benchmark based on the current values of `numRevisions` and `dataSize`.
3. Inside the sub-benchmark, the main loop (`for i := 0; i < b.N; i++`) is run `b.N` times, where `b.N` is determined by the Go testing framework to obtain a reliable benchmark result.
4. In each iteration of the main loop, a new `DoublyLinkedList` is created, and then revisions are generated and inserted into the list. The number of revisions and the data size of each revision are determined by the current values of `numRevisions` and `dataSize`.
5. The `types.NewRandomRevision` function is called with the current index `j` and the specified `dataSize` to create a new revision. Then, a new `DLLNode` is created with the generated revision.
6. The newly created node is inserted into the list at position `j+1` using the `InsertAtPosition` method. This operation is performed `numRevisions` times, resulting in a doubly linked list with the desired number of revisions and data size.

The Go `testing` package measures the time taken for each sub-benchmark and reports the results.

4.1.3 Binary Search Tree

With the implementation of the `Revision` interface and the `Doubly Linked List` data structure complete, the next step was to implement a `Binary Search Tree` (BST) data structure that could also store revisions. The decision to include a BST was primarily driven by its efficient search, insertion, and deletion operations, as well as its ability to

maintain a sorted order of revisions. Furthermore, its hierarchical structure provides an interesting point of comparison against the linear nature of the doubly linked list.

The **Binary Search Tree** implementation consists of a single primary component: the **TreeNode** struct, which represents a node in the tree. In addition, the **TreeNode** struct contains a **Revision** object along with pointers to the left and right child nodes. This structure is designed to efficiently maintain the order of revisions based on their IDs, ensuring that the left subtree contains revisions with smaller IDs and the right subtree contains revisions with larger IDs.

```
type TreeNode struct {
    Revision types.Revision
    Left     *TreeNode
    Right    *TreeNode
}
```

Several methods were implemented to facilitate various operations on the BST, such as inserting revisions, searching for revisions with a specific ID, and removing revisions. These methods provide the necessary functionality to effectively manage the BST and allow for a comprehensive evaluation of its performance characteristics under different scenarios.

```
// Insert inserts a revision into the binary search tree.
func (node *TreeNode) Insert(revision types.Revision) *TreeNode

// Contains checks if a revision with the given ID exists in the binary
// search tree.
func (node *TreeNode) Contains(id int) bool

// Remove removes a revision with the given ID from the binary search tree.
func (node *TreeNode) Remove(id int) (*TreeNode, error)
```

The BST methods were designed to handle the **Revision** interface to maintain consistency with the previous data structures, allowing for seamless integration and comparison between the different data structures. This approach not only simplifies the evaluation process but also emphasises the modularity and adaptability of the code, as multiple data structures can be easily utilised and tested with a shared **Revision** interface.

Tests

TestTreeNode_Insert This test checks if the **Insert** method is working correctly by inserting two revisions into the tree and then asserting that the revisions have been added in the correct order.

TestTreeNode_Contains This test checks if the **Contains** method is working correctly by inserting two revisions into the tree and then verifying if the tree contains those revisions and a non-existent revision.

TestTreeNode_Remove This test checks if the **Remove** method is working correctly by inserting two revisions into the tree, removing one, and then asserting that the removed revision is not in the tree anymore.

Benchmarks

generateRandomRevision - This is a helper function that generates a random revision with a random ID and data size. This function is used in the benchmark tests.

BenchmarkTreeNode_Insert This benchmark measures the performance of the **Insert** method by repeatedly inserting random revisions into the tree.

BenchmarkTreeNode_Contains This benchmark measures the performance of the **Contains** method by first inserting random revisions into the tree and then searching for random revision IDs.

BenchmarkTreeNode_Remove This benchmark measures the performance of the **Remove** method by first inserting random revisions into the tree and then attempting to remove random revision IDs.

BenchmarkTreeNode_RevisionConstruction This benchmark case measures the performance of constructing a binary search tree (BST) with different numbers of revisions and data sizes. The purpose of this benchmark is to evaluate how the BST implementation performs under various conditions and to identify potential bottlenecks or scalability issues.

```
const (  
    smallNumRevisions = 200  
    mediumNumRevisions = 2000
```

```

    largeNumRevisions = 20000
    largeDataSize      = 10000 // 10 Kb
    mediumDataSize     = 1000  // 1 Kb
    smallDataSize      = 100   // 100 bytes
)

func BenchmarkTreeNode_RevisionConstruction(b *testing.B) {
    for _, numRevisions := range []int{smallNumRevisions,
        mediumNumRevisions, largeNumRevisions} {
        for _, dataSize := range []int{smallDataSize, mediumDataSize,
            largeDataSize} {
            b.Run(fmt.Sprintf("numRevisions=%d,dataSize=%d", numRevisions,
                dataSize), func(b *testing.B) {
                for i := 0; i < b.N; i++ {
                    // Create a root node
                    rootRevision := types.NewRandomRevision(0, dataSize)
                    root := &TreeNode{Revision: rootRevision}

                    // Insert revisions
                    for j := 1; j < numRevisions; j++ {
                        revision := types.NewRandomRevision(j, dataSize)
                        root.Insert(revision)
                    }
                }
            })
        }
    }
}

```

The benchmark has an outer loop that iterates over the numbers of revisions and an inner loop that iterates over the data sizes. For each combination of the number of revisions and data size, it runs a sub-benchmark with a unique name in the format "numRevisions=X,dataSize=Y". This makes it easier to analyse the benchmark results and compare the performance across different combinations.

In each sub-benchmark, the following steps are performed:

1. Create a random root revision with ID 0 and the specified data size.
2. Create a new `TreeNode` with the root revision.
3. Insert the remaining revisions into the tree one by one, with unique IDs and the specified data size. The number of inserted revisions is determined by the

current value of the number of revisions.

The Go `testing` package measures the time taken for each sub-benchmark and reports the results.

4.1.4 Directed Acyclic Graph

The next step was to implement a more complex data structure: a **Directed Acyclic Graph** (DAG). This data structure was chosen for its ability to model complex relationships between revisions and represent branching and merging operations commonly found in Version Control Systems (VCS). Additionally, its unique properties and structure provide an interesting point of comparison against the previously implemented data structures.

The **Directed Acyclic Graph** implementation comprises two primary components: the `DAGNode` struct, which represents a node in the graph, and the `DAG struct`, which represents the graph itself. The `DAGNode` struct contains a `Revision` object and two lists representing the parent and child nodes in the graph. In addition, the `DAG` struct maintains a map of nodes indexed by their revision IDs, allowing for efficient retrieval and manipulation of nodes in the graph.

```
type DAGNode struct {
    Revision types.Revision
    Parents []*DAGNode
    Children []*DAGNode
}

type DAG struct {
    nodes map[int]*DAGNode
}
```

A set of methods was developed to facilitate various operations on the DAG, such as adding and removing nodes, adding and removing directed edges between nodes, checking for the existence of nodes with a specific revision ID, and retrieving nodes from the graph. These methods allow for the efficient management of the DAG and enable the evaluation of its performance characteristics under different scenarios.

```
// AddNode adds a new node with the specified revision to the DAG.
func (dag *DAG) AddNode(revision types.Revision) (*DAGNode, error)

// RemoveNode removes a node with the specified revision ID from the DAG.
func (dag *DAG) RemoveNode(revisionID int) error
```

```
// AddEdge adds a directed edge between two nodes with the specified
// revision IDs in the DAG.
func (dag *DAG) AddEdge(parentID, childID int) error

// RemoveEdge removes a directed edge between two nodes with the specified
// revision IDs in the DAG.
func (dag *DAG) RemoveEdge(parentID, childID int) error

// NodeExists checks if a node with the given revision ID exists in the
// DAG.
func (dag *DAG) NodeExists(revisionID int) bool

// GetNode retrieves a node with the given revision ID from the DAG,
// returning nil if it doesn't exist.
func (dag *DAG) GetNode(revisionID int) *DAGNode
```

Tests

TestDAG_AddNode This test checks the `AddNode` function in the DAG. It creates a new DAG and a blank revision with ID 1, then adds the node containing that revision to the DAG. The test asserts that there are no errors and that the revision ID is present in the DAG.

TestDAG_RemoveNode This test checks the `RemoveNode` function in the DAG. It creates a new DAG, adds a node with revision ID 1, and then removes that node. The test asserts that there are no errors and that the revision ID is not present in the DAG after removal.

TestDAG_AddEdge This test checks the `AddEdge` function in the DAG. It creates a new DAG, adds two nodes with revision IDs 1 and 2, and then adds an edge between them. The test asserts that there are no errors and that the edge has been added properly by checking the parent and child connections.

TestDAG_RemoveEdge This test checks the `RemoveEdge` function in the DAG. It creates a new DAG, adds two nodes with revision IDs 1 and 2, adds an edge between them, and then removes that edge. The test asserts that there are no errors and that the edge is removed properly by checking the parent and child connections.

Benchmarks

BenchmarkDAG_AddNode, **BenchmarkDAG_RemoveNode**, **BenchmarkDAG_AddEdge**, and **BenchmarkDAG_RemoveEdge** are benchmark functions that measure the performance of the respective DAG operations (adding nodes, removing nodes, adding edges, and removing edges). They execute the respective functions multiple times with random revisions.

BenchmarkDAG_RevisionConstruction This benchmark case measures the performance of constructing a Directed Acyclic Graph (DAG) with different numbers of nodes and data sizes for the revisions. The purpose of this benchmark is to evaluate the DAG implementation's performance under various conditions and identify potential bottlenecks or scalability issues.

```
const (
    smallNumNodes = 200
    mediumNumNodes = 2000
    largeNumNodes = 20000
    largeDataSize = 10000 // 10 Kb
    mediumDataSize = 1000 // 1 Kb
    smallDataSize = 100    // 100 bytes
)

func BenchmarkDAG_RevisionConstruction(b *testing.B) {
    for _, numNodes := range []int{smallNumNodes, mediumNumNodes,
        largeNumNodes} {
        for _, dataSize := range []int{smallDataSize, mediumDataSize,
            largeDataSize} {
            b.Run(fmt.Sprintf("numRevisions=%d,dataSize=%d", numNodes,
                dataSize), func(b *testing.B) {
                for i := 0; i < b.N; i++ {
                    // Create a DAG
                    dag := NewDAG()
                    // Add nodes with revisions
                    for j := 0; j < numNodes; j++ {
                        revision := types.NewRandomRevision(j, dataSize)
                        dag.AddNode(revision)
                    }

                    // Add edges
                    for j := 0; j < numNodes-1; j++ {
                        dag.AddEdge(j, j+1)
                    }
                }
            })
        }
    }
}
```

For each combination, the benchmark function follows these steps:

1. Create a new DAG using the `NewDAG()` function.
2. Add nodes to the DAG with random revisions. The number of nodes added depends on the current `numNodes` value, and the size of the data in each revision depends on the current `dataSize` value. This is done using the `types.NewRandomRevision()` function and the `AddNode()` method of the DAG.
3. Add edges between the nodes in the DAG. This is done using the `AddEdge()` method of the DAG. In this benchmark, the edges are added sequentially between adjacent nodes (e.g., between nodes 0 and 1, nodes 1 and 2, nodes 2 and 3, and so on).

The benchmark measures the time it takes to perform these operations for each combination of node count and data size.

4.2 Search Algorithms

4.2.1 Linear Search

With the implementation of various data structures complete, the next step was to implement search algorithms that could be used to find revisions within these data structures. The first algorithm implemented was **Linear Search**, a straightforward search method that iterates through the elements of a data structure sequentially, comparing each element to the target value. This algorithm was chosen for its simplicity and as a baseline for comparison with other, more sophisticated search algorithms.

```
func LinearSearch(list *LinkedList.DoublyLinkedList, revisionID int)
    *LinkedList.DLLNode {
currentNode := list.Head
for currentNode != nil {
    if currentNode.Revision.ID() == revisionID {
        return currentNode
    }
}
```



```
    }
    currentNode = currentNode.Next
}
return nil
}
```

The `Linear Search` algorithm was implemented as a function named `LinearSearch`, which takes two parameters: a pointer to a `DoublyLinkedList` object, which represents the list to be searched, and an integer, `revisionID`, representing the revision ID to be searched for in the list. This function was specifically designed to work with the doubly linked list data structure, but it could be easily adapted to work with other data structures as well.

Inside the `LinearSearch` function, a variable named `currentNode` is initialised to the head of the input list. This variable is used to traverse the list, with the traversal continuing until either the target revision ID is found or the end of the list is reached (indicated by a `nil` value for `currentNode`).

A `for` loop is used to traverse the list, with the loop condition checking whether `currentNode` is not `nil`. Inside the loop, an `if` statement compares the revision ID of the `currentNode` to the target `revisionID`. If a match is found, the function returns the `currentNode`. If the end of the list is reached without finding a match, the function returns `nil` to indicate that the target revision ID was not found in the list.

Tests

createTestList - This is a helper function that creates a doubly linked list with 10 nodes, where each node contains a blank revision with IDs from 0 to 9.

TestLinearSearch This test checks if the `LinearSearch` function works as expected. It searches for revisions with IDs from 0 to 9 in the test list and asserts that the nodes are found and have the correct IDs. Additionally, it searches for a revision with an ID of -1, which should not be found, and asserts that the returned node is `nil`.

Benchmarks

BenchmarkLinearSearch_RevisionSearch This benchmark case measures the performance of the `LinearSearch` function on doubly linked lists containing different numbers of revisions and varying revision data sizes. The goal is to understand how the search function performs under various scenarios, which can help identify performance bottlenecks or areas for optimisation.

```

const (
    smallNumRevisions = 200
    mediumNumRevisions = 2000
    largeNumRevisions = 20000
    largeDataSize      = 10000 // 10 Kb
    mediumDataSize     = 1000  // 1 Kb
    smallDataSize      = 100   // 100 bytes
)

func BenchmarkLinearSearch_RevisionSearch(b *testing.B) {
    for _, numRevisions := range []int{smallNumRevisions,
        mediumNumRevisions, largeNumRevisions} {
        for _, dataSize := range []int{smallDataSize, mediumDataSize,
            largeDataSize} {
            b.Run(fmt.Sprintf("numRevisions=%d,dataSize=%d", numRevisions,
                dataSize), func(b *testing.B) {
                list := createListWithRevisions(numRevisions, dataSize)
                b.ResetTimer()

                for i := 0; i < b.N; i++ {
                    revisionID := rand.Intn(numRevisions)
                    LinearSearch(list, revisionID)
                }
            })
        }
    }
}

```

1. The function iterates over different numbers of revisions (small, medium, and large) and data sizes (small, medium, and large) using nested loops. This creates 9 different combinations of revision counts and data sizes for benchmarking.
2. For each revision count and data size combination, the benchmark function executes a sub-benchmark using the `b.Run` method. It constructs a unique name for the sub-benchmark by concatenating the number of revisions and data size (e.g., "numRevisions=200,dataSize=100").
3. Inside the sub-benchmark, the `createListWithRevisions` function creates a doubly linked list with the specified number of revisions and data sizes. Each revision in the list has a unique ID and random data of the specified size.

4. The timer for the benchmark is reset using `b.ResetTimer()`. This ensures that the time taken to create the list is not included in the benchmark results.
5. For each iteration of the benchmark, a random revision ID is generated within the range of the number of revisions in the list. Then, the `LinearSearch` function is called to search for the node containing the revision with the generated ID in the list.
6. The benchmark measures the time taken to perform the linear search for each revision count and data size combination. By analysing the results, you can understand how the performance of the `LinearSearch` function is affected by the number of revisions in the list and the size of the revision data.

4.2.2 Binary Search

The next search algorithm implemented was `Binary Search`, an efficient search method that takes advantage of the sorted order of a data structure, dividing the search interval in half with each iteration. This algorithm was chosen for its efficiency and ability to handle larger datasets effectively, providing a performance comparison against the `Linear Search` algorithm.

```
func BinarySearch(tree *binaryTree.TreeNode, id int) (types.Revision,
    error) {
    currentNode := tree
    for currentNode != nil {
        if id < currentNode.Revision.ID() {
            currentNode = currentNode.Left
        } else if id > currentNode.Revision.ID() {
            currentNode = currentNode.Right
        } else {
            return currentNode.Revision, nil
        }
    }
    return nil, errors.New("revision not found")
}
```

The `Binary Search` algorithm was implemented as a function named `BinarySearch`, which takes two parameters: a pointer to a `TreeNode` object, which represents the root of the binary search tree to be searched, and an integer, `revisionID`, representing the revision ID to be searched for in the tree. This function was specifically designed to work with the binary search tree data structure, leveraging the sorted order of revisions based on their IDs.

Inside the `BinarySearch` function, a variable named `currentNode` is initialised to the root of the input binary tree. This variable is used to traverse the tree, with the traversal continuing until either the target revision ID is found or the end of the tree is reached (indicated by a `nil` value for `currentNode`).

A `for` loop is used to traverse the binary tree. The loop continues until either the target revision ID is found or the end of the tree is reached. Inside the loop, there are three conditions:

- If the target `revisionID` is less than the revision ID of the `currentNode`, the function sets the `currentNode` to the left child of the current node, moving the search down the left subtree.
- If the target `revisionID` is greater than the revision ID of the `currentNode`, the function sets the `currentNode` to the right child of the current node, moving the search down the right subtree.
- If the target `revisionID` is equal to the revision ID of the `currentNode`, the function returns the `Revision` object associated with the `currentNode`.

If the traversal reaches the end of the tree without finding the target revision ID, the function returns a `nil` value to indicate that the target revision ID was not found in the tree.

Tests

createTreeWithRevisions - This helper function creates a binary search tree with a given number of revisions (`numRevisions`) and a specified data size (`dataSize`). It initialises the root node with a random revision and then inserts subsequent revisions into the tree.

TestBinarySearch This test checks if the `BinarySearch` function works as expected. It creates a binary search tree with 100 revisions and a data size of 10 bytes. The test checks for two scenarios:

- (a) The revision with ID 42 is found in the tree. The test asserts that there is no error, the found revision is not `nil`, and its ID matches the expected value (42).
- (b) The revision with ID 999 is not found in the tree. Therefore, the test asserts that there is an error, the error message matches "revision not found", and the returned revision is `nil`.

Benchmarks

BenchmarkBinarySearch_RevisionSearch This benchmark case measures the performance of the `BinarySearch` function under various conditions. The goal is to measure how the search function performs when searching for revisions with different numbers of revisions and varying data sizes in a binary search tree.

```
const (
    smallNumRevisions = 200
    mediumNumRevisions = 2000
    largeNumRevisions = 20000
    largeDataSize      = 10000 // 10 Kb
    mediumDataSize     = 1000  // 1 Kb
    smallDataSize      = 100   // 100 bytes
)

func BenchmarkBinarySearch_RevisionSearch(b *testing.B) {
    for _, numRevisions := range []int{smallNumRevisions,
        mediumNumRevisions, largeNumRevisions} {
        for _, dataSize := range []int{smallDataSize, mediumDataSize,
            largeDataSize} {
            b.Run(fmt.Sprintf("numRevisions=%d,dataSize=%d", numRevisions,
                dataSize), func(b *testing.B) {
                tree := createTreeWithRevisions(numRevisions, dataSize)
                b.ResetTimer()

                for i := 0; i < b.N; i++ {
                    revisionID := rand.Intn(numRevisions)
                    _, _ = BinarySearch(tree, revisionID)
                }
            })
        }
    }
}
```

1. The function iterates over different numbers of revisions (small, medium, and large) and data sizes (small, medium, and large) using nested loops. This creates 9 different combinations of revision counts and data sizes for benchmarking.
2. For each revision count and data size combination, execute a sub-benchmark with the `b.Run` method. This allows each combination to be benchmarked independently, producing individual results. The sub-benchmark is named using a

formatted string containing the number of revisions and data size, e.g., "num-Revisions=200,dataSize=100".

3. Inside each sub-benchmark:

- (a) Create a binary search tree using the `createTreeWithRevisions` function. The tree is populated with the specified number of revisions, each having the specified data size.
- (b) Reset the timer using `b.ResetTimer()`. This ensures that the time spent creating the binary search tree is not included in the benchmark measurement.
- (c) Execute the binary search multiple times, as specified by the `b.N` value, which is determined by the benchmark framework based on the performance of the tested system. In each iteration, a random revision ID within the range of the number of revisions is generated using `rand.Intn(numRevisions)`. Then, the `BinarySearch` function is called to search for the revision with the generated ID in the tree.

The benchmark measures the time taken to perform the binary search for each revision count and data size combination.

4.2.3 Depth-First Search

After implementing both the Linear and Binary search algorithms, the next algorithm to be implemented was **Depth First Search (DFS)**. This graph traversal algorithm explores as far as possible along each branch before backtracking. The primary reason for choosing DFS was its applicability to graph-based data structures, like the **Directed Acyclic Graph (DAG)**.

```
func DepthFirstSearch(dag *graph.DAG, startRevisionID int, visitFunc
    func(*graph.DAGNode)) {
    startNode := dag.GetNode(startRevisionID)
    if startNode == nil {
        return
    }

    visitedNodes := make(map[int]bool)
    depthFirstTraversal(startNode, visitedNodes, visitFunc)
}
```

The Depth First Search algorithm was implemented as a function named `DepthFirstSearch`, which takes three parameters: a pointer to a `graph.DAG` object, which represents the Directed Acyclic Graph to be searched; an integer, `startRevisionID`, representing the revision ID of the starting node for the search; and a function, `visitFunc`, which takes a pointer to a `graph.DAGNode` object as a parameter and is applied to each visited node during the search. The `DepthFirstSearch` function was designed to work with the DAG data structure and can easily handle complex node relationships.

Inside the `DepthFirstSearch` function, the starting node is retrieved from the DAG using the `GetNode` method and the provided `startRevisionID`. If the starting node is `nil`, the function returns immediately, as there is no valid starting point for the search.

Next, the function initialises a map called `visitedNodes` with keys of type `int` (representing revision IDs) and values of type `bool` (indicating whether a node has been visited). This map is used to keep track of visited nodes during the traversal to prevent visiting the same node multiple times.

```
func depthFirstTraversal(node *graph.DAGNode, visitedNodes map[int]bool,
    visitFunc func(*graph.DAGNode)) {
    if visitedNodes[node.Revision.ID()] {
        return
    }

    visitedNodes[node.Revision.ID()] = true
    visitFunc(node)

    for _, childNode := range node.Children {
        depthFirstTraversal(childNode, visitedNodes, visitFunc)
    }
}
```

The `DepthFirstSearch` function then calls a recursive helper function named `depthFirstTraversal`, which takes three parameters: a pointer to a `graph.DAGNode` object representing the current node in the traversal, the `visitedNodes` map, and the `visitFunc`. The `depthFirstTraversal` function is responsible for the actual traversal and visiting of the nodes in the graph.

The `depthFirstTraversal` function starts by checking whether the current node has already been visited by looking up its revision ID in the `visitedNodes` map. If the node has been visited, the function returns immediately to avoid redundant processing. Otherwise, it marks the node as visited by setting the value of the node's revision ID key in the `visitedNodes` map to `true` and applies the `visitFunc` to the

current node.

Finally, the `depthFirstTraversal` function iterates over the `Children` field of the current node, which is a slice of pointers to `graph.DAGNode` objects representing the child nodes. For each child node, the function recursively calls the `depthFirstTraversal` function with the child node, the `visitedNodes` map, and the `visitFunc`.

Tests

createTestDAG - This helper function creates a simple test DAG with five nodes and specific edges between them.

TestDepthFirstSearch This test validates the correctness of the `DepthFirstSearch` function. It creates a test DAG using `createTestDAG` and then performs a depth-first search starting from the node with ID 1. The test function checks if the order of the visited nodes matches the expected order.

Benchmarks

BenchmarkDepthFirstSearch_RevisionSearch This benchmark case measures the performance of the `DepthFirstSearch` function. The goal is to evaluate how the function behaves under various conditions, specifically with different numbers of revisions and data sizes.

```
const (
    smallNumRevisions = 200
    mediumNumRevisions = 2000
    largeNumRevisions = 20000
    largeDataSize      = 10000 // 10 Kb
    mediumDataSize     = 1000  // 1 Kb
    smallDataSize      = 100   // 100 bytes
)

func BenchmarkDepthFirstSearch_RevisionSearch(b *testing.B) {
    for _, numRevisions := range []int{smallNumRevisions,
        mediumNumRevisions, largeNumRevisions} {
        for _, dataSize := range []int{smallDataSize, mediumDataSize,
            largeDataSize} {
            b.Run(fmt.Sprintf("numRevisions=%d,dataSize=%d", numRevisions,
                dataSize), func(b *testing.B) {
                dag := graph.NewDAG()
```



```

nodes := make([]*graph.DAGNode, numRevisions)
for i := 0; i < numRevisions; i++ {
    nodes[i], _ = dag.AddNode(types.NewRandomRevision(i,
        dataSize))
}

for i := 0; i < numRevisions-1; i++ {
    _ = dag.AddEdge(i, i+1)
}

visitor := func(node *graph.DAGNode) {}

b.ResetTimer()
for i := 0; i < b.N; i++ {
    DepthFirstSearch(dag, rand.Intn(numRevisions), visitor)
}
})
}
}
}

```

For each combination of `numRevisions` and `dataSize`, a sub-benchmark is created and executed using the `b.Run` method. Within each sub-benchmark, the following steps are performed:

1. A new DAG is created, and nodes are added with random revisions based on the specified `numRevisions` and `dataSize`. The random revisions are created using the `types.NewRandomRevision` function.
2. An empty visitor function (`visitor := func(node *graph.DAGNode) {}`) is created. This function does nothing when called, ensuring that the benchmark measures only the performance of the DFS algorithm and not any additional processing that might be done in the visit function.
3. The benchmark timer is reset with `b.ResetTimer()`, and the `DepthFirstSearch` function is called `b.N` times using random revision IDs as the starting point.

The benchmark measures the time taken to perform the DFS traversal for each combination of `numRevisions` and `dataSize`.

4.2.4 Breadth-First Search

The final search algorithm implemented was **Breadth First Search (BFS)**. This graph traversal algorithm explores all the nodes at the present depth level before moving on to nodes at the next depth level. BFS was chosen for its applicability to graph-based data structures, like the **Directed Acyclic Graph (DAG)**, and for comparing performance against the **Depth First Search** algorithm.

```
func BreadthFirstSearch(dag *graph.DAG, startRevisionID int, visitFunc
    func(node *graph.DAGNode)) {
    startNode := dag.GetNode(startRevisionID)
    if startNode == nil {
        return
    }

    visitedNodes := make(map[int]bool)
    nodeQueue := []*graph.DAGNode{startNode}

    for len(nodeQueue) > 0 {
        currentNode := nodeQueue[0]
        nodeQueue = nodeQueue[1:]

        if visitedNodes[currentNode.Revision.ID()] {
            continue
        }

        visitFunc(currentNode)
        visitedNodes[currentNode.Revision.ID()] = true

        for _, childNode := range currentNode.Children {
            if !visitedNodes[childNode.Revision.ID()] {
                nodeQueue = append(nodeQueue, childNode)
            }
        }
    }
}
```

The Breadth First Search algorithm was implemented as a function named **BreadthFirstSearch**, which takes three parameters: a pointer to a **graph.DAG** object, representing the **Directed Acyclic Graph** to be searched; an integer, **startRevisionID**, representing the revision ID of the starting node for the search; and a function, **visitFunc**, which takes a pointer to a **graph.DAGNode** object as a parameter and is applied to each vis-

ited node during the search. The `BreadthFirstSearch` function was designed to work with the DAG data structure and can handle complex node relationships efficiently.

Inside the `BreadthFirstSearch` function, the starting node is retrieved from the DAG using the `GetNode` method and the provided `startRevisionID`. If the starting node is `nil`, the function returns immediately, as there is no valid starting point for the search.

A map named `visitedNodes` is initialised to keep track of visited nodes during the search. This map has keys of type `int` (representing revision IDs) and values of type `bool` (indicating whether a node has been visited).

A slice named `nodeQueue` is initialised with the `startNode` as its first element. This queue is used to store nodes that need to be visited during the search, ensuring that nodes are processed in the order of their distance from the starting node.

Tests

TestBreadthFirstSearch This test validates the correctness of the `BreadthFirstSearch` function. The test creates a DAG with 10 nodes, each containing a blank revision with an ID from 0 to 9. Edges are added between consecutive nodes to form a linear graph. The test defines a `visitor` function that appends the revision ID of visited nodes to a `visitedNodes` slice. The test calls `BreadthFirstSearch` with the starting node having an ID of 0 and the defined `visitor` function. The test then checks if the order of visited nodes is as expected.

Benchmarks

BenchmarkBreadthFirstSearch_RevisionSearch This benchmark case measures the performance of the `BreadthFirstSearch` function. It evaluates how the function performs with different numbers of revisions and different data sizes.

```
const (
    smallNumRevisions = 200
    mediumNumRevisions = 2000
    largeNumRevisions = 20000
    largeDataSize     = 10000 // 10 Kb
    mediumDataSize    = 1000  // 1 Kb
    smallDataSize     = 100   // 100 bytes
)

func BenchmarkBreadthFirstSearch_RevisionSearch(b *testing.B) {
    for _, numRevisions := range []int{smallNumRevisions,
        mediumNumRevisions, largeNumRevisions} {
```

```

for _, dataSize := range []int{smallDataSize, mediumDataSize,
    largeDataSize} {
    b.Run(fmt.Sprintf("numRevisions=%d,dataSize=%d", numRevisions,
        dataSize), func(b *testing.B) {
        dag := graph.NewDAG()

        nodes := make([]*graph.DAGNode, numRevisions)
        for i := 0; i < numRevisions; i++ {
            nodes[i], _ = dag.AddNode(types.NewRandomRevision(i,
                dataSize))
        }

        for i := 0; i < numRevisions-1; i++ {
            _ = dag.AddEdge(i, i+1)
        }

        visitor := func(node *graph.DAGNode) {}

        b.ResetTimer()
        for i := 0; i < b.N; i++ {
            BreadthFirstSearch(dag, rand.Intn(numRevisions), visitor)
        }
    })
}
}

```

For each combination of `numRevisions` and `dataSize`, a sub-benchmark is created and executed using the `b.Run` method. Within each sub-benchmark, the following steps are performed:

1. A new DAG is created, and nodes are added with random revisions based on the specified `numRevisions` and `dataSize`. The random revisions are created using the `types.NewRandomRevision` function.
2. An empty visitor function (`visitor := func(node *graph.DAGNode) {}`) is created. This function does nothing when called, ensuring that the benchmark measures only the performance of the BFS algorithm and not any additional processing that might be done in the visit function.
3. The benchmark timer is reset with `b.ResetTimer()`, and the `BreadthFirstSearch` function is called `b.N` times using random revision IDs as the starting point.

The benchmark measures the time taken to perform the BFS traversal for each combination of `numRevisions` and `dataSize`.

4.3 Dynamic Programming Algorithms

4.3.1 Longest Common Subsequence

Inefficient Algorithm

Efficient Algorithm

Tests

Benchmarks

4.3.2 Longest Increasing Subsequence

Inefficient Algorithm

Efficient Algorithm

Tests

Benchmarks

4.4 Diffing Algorithms

4.4.1 Myers Diff

Tests

Benchmarks

4.4.2 Patience Diff

Tests

Benchmarks

Chapter 5

Evaluation

Chapter 6

Conclusion

List of Figures

2.1	Timeline of the Creation of Version Control Systems [9]	4
2.2	Local Version Control System (LVCS)	5
2.3	Centralized Version Control System (CVCS)	9
2.4	Distributed Version Control System (DVCS)	13
3.1	Singly Linked List (SLL) [12]	18
3.2	Doubly Linked List (DLL) [12]	19
3.3	Binary Tree [2]	20
3.4	Hash Table [8]	23
3.5	Directed Acyclic Graph (DAG) [10]	25
3.6	Linear Search Algorithm [6]	27
3.7	Binary Search Algorithm [3]	28
3.8	Depth First Search Algorithm [13]	30
3.9	Breadth First Search Algorithm [13]	31

List of Tables

3.1	Efficiency Analysis of Singly-Linked List Operations	17
3.2	Efficiency Analysis of Doubly-Linked List Operations	18
3.3	Efficiency Analysis of Binary Tree Operations	20
3.4	Efficiency Analysis of Hash Table Operations	22

Bibliography

- [1] J. Cloer. “10 years of git: An interview with git creator linus torvalds.” (Aug. 2019), [Online]. Available: <https://www.linux.com/news/10-years-git-interview-git-creator-linus-torvalds/>.
- [2] C. McMahon. “Understanding binary search trees.” (Jul. 2020), [Online]. Available: <https://dev.to/christinamcmahon/understanding-binary-search-trees-4d90>.
- [3] J. P. “How to do a binary search in python.” (Nov. 2020), [Online]. Available: <https://learncodingfast.com/binary-search-in-python/>.
- [4] O. Popovi. “Binary search in python.” (Jun. 2020), [Online]. Available: <https://stackabuse.com/binary-search-in-python/>.
- [5] A. S. Ravikiran. “Linked list in a data structure: All you need to know.” (Nov. 2022), [Online]. Available: <https://www.simplilearn.com/tutorials/data-structure-tutorial/linked-list-in-data-structure>.
- [6] A. S. Ravikiran. “What is linear search algorithm?” (Feb. 2023), [Online]. Available: <https://www.simplilearn.com/tutorials/data-structure-tutorial/linear-search-algorithm>.
- [7] Shaalaa. “Discuss linear search algorithm.” (), [Online]. Available: https://www.shaalaa.com/question-bank-solutions/discuss-linear-search-algorithm-algorithm-for-searching-techniques_229584.
- [8] K. Stemmler. “Hash tables: What, why, how to use them.” (Jan. 2022), [Online]. Available: <https://khalilstemmler.com/blogs/data-structures-algorithms/hash-tables/>.
- [9] J. Stopak. “Version control systems: A technical guide to vcs internals.” (Nov. 2019), [Online]. Available: <https://initialcommit.com/blog/Technical-Guide-VCS-Internals>.
- [10] H. Surti. “Advanced data structures part 1: Directed acyclic graph (dag).” (Apr. 2016), [Online]. Available: <https://medium.com/@hamzasurti/advanced-data-structures-part-1-directed-acyclic-graph-dag-c1d1145b5e5a>.

- [11] W. F. Tichy, "Rcs - a system for version control," *Software: Practice and Experience*, vol. 15, no. 7, pp. 637–654, 1985. DOI: 10.1002/spe.4380150703.
- [12] R. Vaghani. "Difference between singly linked list and doubly linked list." (Jan. 2023), [Online]. Available: <https://www.geeksforgeeks.org/difference-between-singly-linked-list-and-doubly-linked-list/>.
- [13] D. Zaltsman. "Difference between depth first search and breadth first search." (Mar. 2020), [Online]. Available: <https://dev.to/danimal92/difference-between-depth-first-search-and-breadth-first-search-6om>.